

# TODO: Exercise 6-2, 6-3, 8-1, 8-2, 10-1

## Import Pandas

```
In [2]: import pandas as pd
```

## Set up connection to DB

```
In [1]: import os
        from dotenv import find_dotenv, dotenv_values

        keys = list(dotenv_values(find_dotenv('.env')).items())
        os.environ['POSTGRES_PASS'] = keys[1][1]
        os.environ['POSTGRES_USER'] = keys[2][1]
        host = 'localhost'
        port = '5432'
        db = 'bank'

        engine=f'postgresql://{os.getenv('POSTGRES_USER')}:{os.getenv('POSTGRES_PASS')}@{ho
```

## Exercise 6-2

Write a compound query that finds the first and last names of all individual customers along with the first and last names of all employees.

```
In [5]: pd.read_sql_query("""
        SELECT fname, lname
        FROM individual
        UNION
        SELECT fname, lname
        FROM employee
        """,
        engine
    )
```

Out[5]:

	<b>fname</b>	<b>lname</b>
<b>0</b>	Paula	Roberts
<b>1</b>	Susan	Tingley
<b>2</b>	Samantha	Jameson
<b>3</b>	Cindy	Mason
<b>4</b>	Chris	Tucker
<b>5</b>	Richard	Farley
<b>6</b>	Frank	Portman
<b>7</b>	Rick	Tulman
<b>8</b>	Beth	Fowler
<b>9</b>	Sarah	Parker
<b>10</b>	John	Gooding
<b>11</b>	John	Hayward
<b>12</b>	Robert	Tyler
<b>13</b>	Margaret	Young
<b>14</b>	Charles	Frasier
<b>15</b>	Thomas	Ziegler
<b>16</b>	Theresa	Markham
<b>17</b>	Susan	Barker
<b>18</b>	Louis	Blake
<b>19</b>	John	Blake
<b>20</b>	James	Hadley
<b>21</b>	Susan	Hawthorne
<b>22</b>	John	Spencer
<b>23</b>	Jane	Grossman
<b>24</b>	Frank	Tucker
<b>25</b>	Helen	Fleming
<b>26</b>	Michael	Smith

## Exercise 6-3

Sort the results from Exercise 6-2 by the lname column.

```
In [6]: pd.read_sql_query("""
        SELECT fname, lname
        FROM individual
        UNION
        SELECT fname, lname
        FROM employee
        ORDER BY lname ASC
        """,
        engine
    )
```

Out[6]:

	<b>fname</b>	<b>lname</b>
<b>0</b>	Susan	Barker
<b>1</b>	Louis	Blake
<b>2</b>	John	Blake
<b>3</b>	Richard	Farley
<b>4</b>	Helen	Fleming
<b>5</b>	Beth	Fowler
<b>6</b>	Charles	Frasier
<b>7</b>	John	Gooding
<b>8</b>	Jane	Grossman
<b>9</b>	James	Hadley
<b>10</b>	Susan	Hawthorne
<b>11</b>	John	Hayward
<b>12</b>	Samantha	Jameson
<b>13</b>	Theresa	Markham
<b>14</b>	Cindy	Mason
<b>15</b>	Sarah	Parker
<b>16</b>	Frank	Portman
<b>17</b>	Paula	Roberts
<b>18</b>	Michael	Smith
<b>19</b>	John	Spencer
<b>20</b>	Susan	Tingley
<b>21</b>	Frank	Tucker
<b>22</b>	Chris	Tucker
<b>23</b>	Rick	Tulman
<b>24</b>	Robert	Tyler
<b>25</b>	Margaret	Young
<b>26</b>	Thomas	Ziegler

## Exercise 8-1

Construct a query that counts the number of rows in the account table.

```
In [10]: pd.read_sql_query("""
            SELECT COUNT(*) rows
            FROM account
            """,
            engine
        )
```

```
Out[10]:
```

	rows
0	24

## Exercise 8-2

Modify your query from Exercise 8-1 to count the number of accounts held by each customer. Show the customer ID and the number of accounts for each customer.

```
In [22]: pd.read_sql_query("""
            SELECT cust_id,
            COUNT(*) num_accounts
            FROM account
            GROUP BY cust_id
            ORDER BY cust_id ASC
            """,
            engine
        )
```

```
Out[22]:
```

	cust_id	num_accounts
0	1	3
1	2	2
2	3	2
3	4	3
4	5	1
5	6	2
6	7	1
7	8	2
8	9	3
9	10	2
10	11	1
11	12	1
12	13	1

## Exercise 10-1

Write a query that returns all product names along with the accounts based on that product (use the product\_cd column in the account table to link to the product table). Include all products, even if no accounts have been opened for that product.

```
In [44]: """
This will require the following steps:
1. Select account id and product name
2. Full outer join on product_cd... this is required to return product names that a

After completing this, I am noticing that the type of account ID has changed. I thi
Maybe they are represented as floats? Thus, the column is cast to a float to have a
I don't think this matters for this application, but I will fix it anyway. I should

I will need some AI assistance on this, so I will output the data to a csv.
"""

pd.read_sql_query("""
    SELECT a.account_id, p.name
    FROM account a FULL OUTER JOIN product p
    ON a.product_cd = p.product_cd
    """,
    engine
).to_csv("data/10-1.csv")
```

I consulted with Gemini on this problem, conversation link, below:

<https://g.co/gemini/share/726570e92dc5>

I am going to give the SQL statement that Gemini provided a shot.

```
In [49]: pd.read_sql_query("""
    SELECT COALESCE(a.account_id, -1) AS account_id, -- Replace NUL
    p.name
    FROM account a FULL OUTER JOIN product p
    ON a.product_cd = p.product_cd
    """,
    engine
)
```

Out[49]:

	<b>account_id</b>	<b>name</b>
<b>0</b>	1	checking account
<b>1</b>	2	savings account
<b>2</b>	3	certificate of deposit
<b>3</b>	4	checking account
<b>4</b>	5	savings account
<b>5</b>	7	checking account
<b>6</b>	8	money market account
<b>7</b>	10	checking account
<b>8</b>	11	savings account
<b>9</b>	12	money market account
<b>10</b>	13	checking account
<b>11</b>	14	checking account
<b>12</b>	15	certificate of deposit
<b>13</b>	17	certificate of deposit
<b>14</b>	18	checking account
<b>15</b>	19	savings account
<b>16</b>	21	checking account
<b>17</b>	22	money market account
<b>18</b>	23	certificate of deposit
<b>19</b>	24	checking account
<b>20</b>	25	business line of credit
<b>21</b>	27	business line of credit
<b>22</b>	28	checking account
<b>23</b>	29	small business loan
<b>24</b>	-1	home mortgage
<b>25</b>	-1	auto loan

Okay, this is better. I wish I could set -1 to NULL though.

I followed up with Gemini and now I have a better understanding of why this occurs, conversation, below:

<https://g.co/gemini/share/2a47f4b5a45a>

So, Pandas is at fault for the float casting, not SQL. The previous solution just modified the SQL output prior to Pandas interpreting it, so it appeared to do what I want.

Turns out, Gemini was right to suggest I need to modify the dataframe the first time around. Who would have guessed...

I guess I will oblige.

```
In [50]: df = pd.read_sql_query("""
        SELECT a.account_id, p.name
        FROM account a FULL OUTER JOIN product p
        ON a.product_cd = p.product_cd
        """,
        engine
    )
df['account_id'] = df['account_id'].astype('Int64')
df
```



Out[50]:

	<b>account_id</b>	<b>name</b>
<b>0</b>	1	checking account
<b>1</b>	2	savings account
<b>2</b>	3	certificate of deposit
<b>3</b>	4	checking account
<b>4</b>	5	savings account
<b>5</b>	7	checking account
<b>6</b>	8	money market account
<b>7</b>	10	checking account
<b>8</b>	11	savings account
<b>9</b>	12	money market account
<b>10</b>	13	checking account
<b>11</b>	14	checking account
<b>12</b>	15	certificate of deposit
<b>13</b>	17	certificate of deposit
<b>14</b>	18	checking account
<b>15</b>	19	savings account
<b>16</b>	21	checking account
<b>17</b>	22	money market account
<b>18</b>	23	certificate of deposit
<b>19</b>	24	checking account
<b>20</b>	25	business line of credit
<b>21</b>	27	business line of credit
<b>22</b>	28	checking account
<b>23</b>	29	small business loan
<b>24</b>	<NA>	home mortgage
<b>25</b>	<NA>	auto loan

Well, that did it. I apologize for the lengthy side quest. I should have recognized that I was looking at a pandas dataframe representation of the SQL query, therefore it would be a pandas issue.